



- **Reward is negative:** The objective reduces to

$$\max \left(\frac{\pi_{\theta}(a^{(t)} \mid s^{(t)})}{\pi_{\theta_{\text{old}}}(a^{(t)} \mid s^{(t)})}, (1 - \epsilon) \right) R^{(t)}$$

Then, the objective decreases with $\pi_{\theta}(a^{(t)} \mid s^{(t)})$.

Once $\pi_{\theta}(a^{(t)} \mid s^{(t)}) < (1 - \epsilon)\pi_{\theta_{\text{old}}}(a^{(t)} \mid s^{(t)})$,

the max kicks in, with a ceiling of $(1 - \epsilon)R^{(t)}$.

Search as an optimization problem

- **Reward is positive:** The objective reduces to

$$\min \left(\frac{\pi_{\theta}(a^{(t)} \mid s^{(t)})}{\pi_{\theta_{\text{old}}}(a^{(t)} \mid s^{(t)})}, (1 + \epsilon) \right) R^{(t)}$$

Then, the objective increases with $\pi_{\theta}(a^{(t)} \mid s^{(t)})$.

Once $\pi_{\theta}(a^{(t)} \mid s^{(t)}) > (1 + \epsilon)\pi_{\theta_{\text{old}}}(a^{(t)} \mid s^{(t)})$,

the min kicks in, with a ceiling of $(1 + \epsilon)R^{(t)}$.

$$g(\epsilon, R^{(t)}) = \begin{cases} (1 + \epsilon) R^{(t)} & R^{(t)} \geq 0 \\ (1 - \epsilon) R^{(t)} & R^{(t)} < 0 \end{cases}$$

[OpenAI SpinningUp, Proximal Policy Optimization]

PPO-Clip as a Regularizer

- **Reward is negative:** The objective reduces to

$$\max \left(\frac{\pi_{\theta}(a^{(t)} \mid s^{(t)})}{\pi_{\theta_{\text{old}}}(a^{(t)} \mid s^{(t)})}, (1 - \epsilon) \right) R^{(t)}$$

Then, the objective decreases with $\pi_{\theta}(a^{(t)} \mid s^{(t)})$.

Once $\pi_{\theta}(a^{(t)} \mid s^{(t)}) < (1 - \epsilon)\pi_{\theta_{\text{old}}}(a^{(t)} \mid s^{(t)})$,

the max kicks in, with a ceiling of $(1 - \epsilon)R^{(t)}$.

- **Reward is positive:** The objective reduces to

$$\min \left(\frac{\pi_{\theta}(a^{(t)} \mid s^{(t)})}{\pi_{\theta_{\text{old}}}(a^{(t)} \mid s^{(t)})}, (1 + \epsilon) \right) R^{(t)}$$

Then, the objective increases with $\pi_{\theta}(a^{(t)} \mid s^{(t)})$.

Once $\pi_{\theta}(a^{(t)} \mid s^{(t)}) > (1 + \epsilon)\pi_{\theta_{\text{old}}}(a^{(t)} \mid s^{(t)})$,

the min kicks in, with a ceiling of $(1 + \epsilon)R^{(t)}$.